

Choose your own project

KC Cardakli

2024-08-07

Overview

In this project, we look at the US COVID deaths from Jan 2024 - Sept 23, 2023 and associated diagnoses at the time of death. We fit models using the associated diagnoses predicting the number of deaths. Once we find the best models we identify the variable importance for these models.

We attempt to identify the correlated factors in COVID19 deaths.

The input to our models are the number of diagnoses for each row and the output is the COVID19 deaths for that time period and geographic location.

We use dataset from US Health and Human Services (HHS). We first wrangle the data and eliminate duplicate data and rows with missing data. Please refer to the section “About the dataset” for further details.

We then convert the data into wide format and check against the total COVID19 deaths from CDC website. We construct 14 models and train these models via bootstrapping and cross validation with a k value of 25.

The models that provide best RMSE (root of mean square error) values are provided by the xgbTree and rf algorithms followed by gamLoess. xgbTree and rf are Tree-Based Models, xgbTree is Extreme Gradient Boosting Model where as rf is Random Forest Model. They provide similar variable importance. The third best model, gamLoess, combines Generalized Additive Models (GAM) with LOESS (Locally Estimated Scatterplot Smoothing) and gives a different variable importance. Please see below for details.

The neural network model(nnet), provided one the worst models with an RMSE value of 2186. We investigate to see if we can improve on this model by changing the model parameters, number of units in the hidden layer and decay. Although we improve the RMSE value significantly, it is still not as accurate as xgbTree, rf, or gamLoess models. Apparently nnet model can have only one hidden layer, and that may be limitation of this model.

When looking at the results of this project, It is probably important remember that correlation doesn't mean causality; some of the diagnoses could be (and are) complications of COVID19, rather than causes. The best model gives us the 10 most important variables, in the order of decreasing importance, as respiratory failure(COVID19 itself causes respiratory failure), influenza and pneumonia, malignant neoplasms (cancer), adult respiratory distress syndrome, cardiac arrhythmia, respiratory arrest, chronic lower respiratory diseases, diabetes, ischemic heart disease, and heart failure. The result seems reasonable, but of course much more research would have to be done to generalize.

It would be very interesting to get the list of the diagnoses when the patient was diagnosed with COVID19 not at the time of death. And of course, it would be even better if we could get the diagnoses of every COVID19 patient when they got sick, not just the ones that died from it.

About The Dataset

For this project we use a public dataset from US Health and Human Services (HHS):

Conditions Contributing to COVID-19 Deaths, by State and Age, Provisional 2020-2023

<https://catalog.data.gov/dataset/conditions-contributing-to-deaths-involving-coronavirus-disease-2019-covid-19-by-age-group>

The dataset summarizes the COVID19 deaths and associated factors.

I suspect the underlying full dataset is at the US Center for Disease Control (CDC):

<https://wonder.cdc.gov/mcd.html>

but this database is not public, and only available to researchers with certain conditions:

<https://wonder.cdc.gov/mcd-icd10-provisional.html>

Among the conditions is “Do not present or publish death counts of 9 or fewer or death rates based on counts of nine or fewer (in figures, graphs, maps, tables, etc.).”

I have asked the TAs of this course about the CDC dataset and I was asked not use CDC dataset since I could not provide the dataset as part of the project.

HHS dataset puts NA for values 1-9 as CDC dataset conditions dictate.

HHS dataset also provides the national sums.

HHS dataset provides the associated diagnoses for COVID19 deaths in each row, data is tabulated as on a per month / state basis.

In the State column, possible values are all US states, Washington DC, and Puerto Rico.

Methods

We use 14 models:

| | |
|----------|--|
| lm | Linear Regression |
| glm | Generalized Linear Model |
| knn | K-Nearest Neighbors |
| rf | Random Forest |
| gamLoess | Generalized Additive Model (GAM) combined with LOESS (Locally Estimated Scatterplot Smoothing) |
| rpart | Recursive Partitioning and Regression Trees |
| xgbTree | Extreme Gradient Boosting (XGBoost) |
| cforest | Conditional Inference Trees |
| glmnet | Regularized Generalized Linear Models (Elastic Net) |
| bayesglm | Bayesian Generalized Linear Models |
| pcr | Principal Component Regression |
| pls | Partial Least Squares Regression |
| ridge | Ridge Regression |
| nnet | Neural Network |

We first start with bootstrapping k value of 25, and then we use cross validation with a k-value of 25.

Results

Cross validation with a k value of 25 provides the best RMSE values on the test data.

| Model | RMSE |
|---------|----------|
| xgbTree | 161.1663 |

| Model | RMSE |
|----------|----------|
| rf | 248.5894 |
| gamLoess | 252.0772 |

Bootstrapping yields similar results; the order of models do not change:

| Model | RMSE |
|----------|----------|
| xgbTree | 231.0553 |
| rf | 248.5894 |
| gamLoess | 252.0772 |

| xgbTree variable importance gives us the model's correlated diagnoses | |
|---|-----------|
| Respiratory.failure | 100.00000 |
| Influenza.pneumonia | 90.07434 |
| Malignant.neoplasms | 30.88395 |
| Adult.respiratory.distress.syndrome | 21.18099 |
| Cardiac.arrhythmia | 18.13848 |
| Respiratory.arrest | 14.79050 |
| Chronic.lower.respiratory.diseases | 12.53784 |
| Diabetes | 5.39707 |
| Ischemic.heart.disease | 5.32439 |
| Heart.failure | 4.20999 |
| Hypertensive.diseases | 2.35590 |
| Vascular.unspecified.dementia | 1.06199 |
| Cardiac.arrest | 0.32535 |
| Sepsis | 0.29698 |
| Other.diseases.of.the.respiratory.system | 0.22408 |
| Cerebrovascular.diseases | 0.17067 |
| Renal.failure | 0.13228 |
| Obesity | 0.10758 |
| Alzheimer.disease | 0.03441 |
| Other.diseases.of.the.circulatory.system | 0.00000 |

| Variable Importance for rf model is similar: | |
|--|----------|
| Influenza.pneumonia | 100.0000 |
| Respiratory.failure | 84.0869 |
| Ischemic.heart.disease | 67.4816 |
| Cardiac.arrhythmia | 44.0161 |
| Other.diseases.of.the.circulatory.system | 23.0288 |
| Cerebrovascular.diseases | 16.9564 |
| Renal.failure | 16.5231 |
| Diabetes | 13.6520 |
| Adult.respiratory.distress.syndrome | 13.1184 |
| Respiratory.arrest | 9.0165 |
| Chronic.lower.respiratory.diseases | 7.5493 |
| Heart.failure | 6.6621 |
| Hypertensive.diseases | 5.5820 |
| Cardiac.arrest | 4.9089 |

| | |
|--|--------|
| Variable Importance for rf model is similar: | |
| Alzheimer.disease | 3.1151 |
| Malignant.neoplasms | 2.4882 |
| Other.diseases.of.the.respiratory.system | 2.2472 |
| Sepsis | 1.5487 |
| Vascular.unspecified.dementia | 0.9601 |
| Obesity | 0.0000 |

| | |
|---|---------|
| Variable Importance for gamLoess model is completely different: | |
| Other.diseases.of.the.circulatory.system | 100.000 |
| Cardiac.arrest | 84.522 |
| Renal.failure | 83.760 |
| Sepsis | 80.355 |
| Cerebrovascular.diseases | 68.651 |
| Chronic.lower.respiratory.diseases | 57.188 |
| Diabetes | 51.617 |
| Obesity | 49.864 |
| Adult.respiratory.distress.syndrome | 49.663 |
| Respiratory.arrest | 45.461 |
| Hypertensive.diseases | 35.356 |
| Respiratory.failure | 34.682 |
| Other.diseases.of.the.respiratory.system | 27.160 |
| Influenza.pneumonia | 20.816 |
| Alzheimer.disease | 17.822 |
| Cardiac.arrhythmia | 16.083 |
| Heart.failure | 12.874 |
| Malignant.neoplasms | 8.244 |
| Vascular.unspecified.dementia | 3.550 |
| Ischemic.heart.disease | 0.000 |

Conclusions

Tree based models give us the best RMSE results. In particular xgbTree model gives the best results with 25-fold cross validation. Considering the mean of the output parameter, number of COVID19 deaths, is 1828.651, and median is 1236, a RMSE of 161.1663 seems good.

Of course, it may be possible to improve on this RMSE value by using the base source data from CDC.

It was also interesting to see that tree models have similar variable importance, while gamLoess has a much different variable importance

Possible Future Work

It would be interesting to replicate the work with the non-public base data from CDC with the actual counts 1-9 included. That would give us more data rows, and it would be more precise.

Also as mentioned before, it would be interesting to get the diagnoses of the patients when they get contract COVID19, then we can compare the possible associations (possibly risk factors) between the patients that survived COVID19 and not.

References

Public dataset from US Health and Human Services (HHS): Conditions Contributing to COVID-19 Deaths, by State and Age, Provisional 2020-2023

<https://catalog.data.gov/dataset/conditions-contributing-to-deaths-involving-coronavirus-disease-2019-covid-19-by-age-group>

Libraries

```
#####  
#  
# Please uncomment these lines to install packages as needed in your system:  
#  
#install.packages("dplyr")  
#install.packages("data.table")  
#install.packages("readr")  
#install.packages("caret")  
#install.packages("tidyr")  
  
library("dplyr")
```

```
##  
## Attaching package: 'dplyr'  
  
## The following objects are masked from 'package:stats':  
##  
##   filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
library("data.table")
```

```
##  
## Attaching package: 'data.table'  
  
## The following objects are masked from 'package:dplyr':  
##  
##   between, first, last
```

```
library("readr")  
library("caret")
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library("tidyr")

options(max.print = 500)
```

Reading the HHS dataset and checking one month and one state visually

Please don't forget to modify your working directory as needed. The R, Rmd, and pdf files and csv file from HHS should be in your working directory.

```
#####
# !!!!
# IMPORTANT: You may modify the directory to your working directory as needed
#
print("working directory: ")
```

```
## [1] "working directory: "
```

```
# setwd("/Users/kccardakli/Documents/R/projects/Capstone Project/CYO")
getwd()
```

```
## [1] "/Users/kccardakli/Documents/R/projects/Capstone Project/CYO"
```

```
options(timeout = 1200)
```

```
# Specifying the URL of the uncompressed file from CDC
# If this fails for you, you can download the zipped file from GitHub, and unzip it manually:
# https://github.com/KCardakli/CYO/blob/4c86003ad5095c7baceed361753182a35b30a439/Conditions_Contributin
# or the uncompressed file from Google drive:
# https://drive.google.com/file/d/1WQt3y2N-xiEfEiuI4VuExcexpjjchNe2/view?usp=sharing
# R-studio fails when large files are downloaded from GitHub or Google drive:
url <- "https://data.cdc.gov/api/views/hk9y-quqm/rows.csv?accessType=DOWNLOAD"
```

```
# Name of the downloaded file
destfile <- "Conditions_Contributing_to_COVID-19_Deaths__by_State_and_Age__Provisional_2020-2023.csv"
```

```
# Downloading the file
if(!file.exists(destfile))
  download.file(url, destfile, method = "libcurl")
```

```
# reading the HHS dataset
data <- read.csv("./Conditions_Contributing_to_COVID-19_Deaths__by_State_and_Age__Provisional_2020-2023
```

```
# Let's check the data for one month and one state
# We will use the Age.Group == "All Ages" as it minimizes the effect of putting NA for
# values 1-9 as explained above.
#
```

```
temp <- data |> filter(Month == 8 & Year == 2020 & State == "Alabama" & Age.Group == "All Ages") |>
  select(-Data.As.Of, -Start.Date, -End.Date, -ICD10_codes)
write_csv(temp, "temp.csv")
```

Data Wrangling

```
# Eliminate the nationwide data as they replicate the statewide data by summation.
# We don't want to double count the data.
# We will use the monthly data.
# Eliminate the rows that have "" or NA for the COVID.19.Deaths, this comes from the CDC
# source data where "one or more data cells have counts between 1-9 and have been suppressed
# in accordance with NCHS confidentiality standards". This elimination introduces some error,
# solution would be to use the CDC dataset, but it is not public.
# We also eliminate "All other conditions and causes (residual)", "Intentional and unintentional
# injury, poisoning, and other adverse events" as these are catchall phrases and are not specific
# diagnoses.
# We are using the COVID.19.Deaths and not Number.of.Mentions as some conditions are
# mentioned multiple times in death certificates under reason for death and contributing
# factors, so they could be counted more than once.
data <- data |> filter(data$Group == "By Month" &
  data$COVID.19.Deaths != "" &
  State != "United States" &
  !is.na(data$COVID.19.Deaths) &
  Condition != 'All other conditions and causes (residual)' &
  Condition != 'Intentional and unintentional injury, poisoning, and other adverse events' &
  Age.Group == 'All Ages') |> select(Year, Month, State, COVID.19.Deaths, Condition)
str(data)
```

```
## 'data.frame': 35091 obs. of 5 variables:
## $ Year : int 2020 2020 2020 2020 2020 2020 2020 2020 2020 2020 2020 ...
## $ Month : int 1 2 3 4 5 6 7 8 9 10 ...
## $ State : chr "Alabama" "Alabama" "Alabama" "Alabama" ...
## $ COVID.19.Deaths: int 0 0 11 92 128 125 264 318 188 183 ...
## $ Condition : chr "Influenza and pneumonia" "Influenza and pneumonia" "Influenza and pneumoni
```

```
# Let's check the total number of COVID deaths in the dataset
# It is inline with other resources from CDC (1.14M):
# https://www.cdc.gov/nchs/nvss/vsrr/covid19/index.htm
#
t <- data |> filter(Condition == "COVID-19" & !is.na(COVID.19.Deaths))
print("Total COVID deaths in the dataset 01/01/2020 - 09/23/2023:")
```

```
## [1] "Total COVID deaths in the dataset 01/01/2020 - 09/23/2023:"
```

```
sum(t$COVID.19.Deaths)
```

```
## [1] 1152167
```

```
# We now rename the fields that will become columns so that they are easily readable and can be
# accessed without using `` characters
data <- data |>
  mutate(Condition = gsub(",", "", Condition)) |>
  mutate(Condition = gsub(" and ", " ", Condition)) |>
  mutate(Condition = gsub('-', "", Condition)) |>
  mutate(Condition = gsub(" ", " ", Condition)) |>
```

```
mutate(Condition = gsub(" ", ".", Condition))
```

```
#These are the unique conditions
unique(data$Condition)
```

```
## [1] "Influenza.pneumonia"
## [2] "Chronic.lower.respiratory.diseases"
## [3] "Adult.respiratory.distress.syndrome"
## [4] "Respiratory.failure"
## [5] "Respiratory.arrest"
## [6] "Other.diseases.of.the.respiratory.system"
## [7] "Hypertensive.diseases"
## [8] "Ischemic.heart.disease"
## [9] "Cardiac.arrest"
## [10] "Cardiac.arrhythmia"
## [11] "Heart.failure"
## [12] "Cerebrovascular.diseases"
## [13] "Other.diseases.of.the.circulatory.system"
## [14] "Sepsis"
## [15] "Malignant.neoplasms"
## [16] "Diabetes"
## [17] "Obesity"
## [18] "Alzheimer.disease"
## [19] "Vascular.unspecified.dementia"
## [20] "Renal.failure"
## [21] "COVID19"
```

```
# Form the wide data with "Condition"s are columns
wide_data <- data |>
  pivot_wider(id_cols = c(Year, Month, State),
              names_from = Condition,
              values_from = c(COVID.19.Deaths))
str(wide_data)
```

```
## tibble [2,385 x 24] (S3: tbl_df/tbl/data.frame)
## $ Year : int [1:2385] 2020 2020 2020 2020 2020 2020 2020 2020 2020 2020
## $ Month : int [1:2385] 1 2 3 4 5 6 7 8 9 10 ...
## $ State : chr [1:2385] "Alabama" "Alabama" "Alabama" "Alabama" ...
## $ Influenza.pneumonia : int [1:2385] 0 0 11 92 128 125 264 318 188 183 ...
## $ Chronic.lower.respiratory.diseases : int [1:2385] 0 NA NA 36 32 31 73 55 26 30 ...
## $ Adult.respiratory.distress.syndrome : int [1:2385] 0 0 NA 32 24 28 42 67 38 40 ...
## $ Respiratory.failure : int [1:2385] 0 0 21 126 144 134 299 380 234 245 ...
## $ Respiratory.arrest : int [1:2385] NA 0 0 14 39 32 74 71 37 39 ...
## $ Other.diseases.of.the.respiratory.system: int [1:2385] 0 0 NA 15 NA 17 26 34 22 22 ...
## $ Hypertensive.diseases : int [1:2385] NA 0 NA 52 51 35 117 105 55 51 ...
## $ Ischemic.heart.disease : int [1:2385] 0 0 NA 31 39 32 82 82 44 49 ...
## $ Cardiac.arrest : int [1:2385] NA 0 NA 43 60 68 116 128 61 68 ...
## $ Cardiac.arrhythmia : int [1:2385] NA 0 NA 24 27 18 60 58 31 33 ...
## $ Heart.failure : int [1:2385] 0 0 NA 23 36 16 58 55 35 38 ...
## $ Cerebrovascular.diseases : int [1:2385] 0 0 NA 20 22 23 30 41 21 27 ...
## $ Other.diseases.of.the.circulatory.system: int [1:2385] NA NA NA 18 27 15 42 61 40 35 ...
## $ Sepsis : int [1:2385] 0 0 10 51 54 51 87 112 70 68 ...
## $ Malignant.neoplasms : int [1:2385] 0 0 NA 11 12 12 35 12 24 22 ...
```



```
## $ Diabetes : int [1:2385] NA NA NA 39 47 42 95 110 68 58 ...
## $ Obesity : int [1:2385] 0 0 0 NA NA NA 26 30 16 10 ...
## $ Alzheimer.disease : int [1:2385] 0 0 NA 15 18 NA 44 31 16 17 ...
## $ Vascular.unspecified.dementia : int [1:2385] 0 0 NA 31 50 35 68 71 65 39 ...
## $ Renal.failure : int [1:2385] 0 0 NA 28 47 38 96 82 57 48 ...
## $ COVID19 : int [1:2385] NA NA 48 339 430 399 879 964 568 560 ...
```

```
wide_data |> print(, n = 25)
```

```
## # A tibble: 2,385 x 24
##   Year Month State Influenza.pneumonia Chronic.lower.respiratory.diseases
##   <int> <int> <chr> <int> <int>
## 1 2020 1 Alabama 0 0
## 2 2020 2 Alabama 0 NA
## 3 2020 3 Alabama 11 NA
## 4 2020 4 Alabama 92 36
## 5 2020 5 Alabama 128 32
## 6 2020 6 Alabama 125 31
## 7 2020 7 Alabama 264 73
## 8 2020 8 Alabama 318 55
## 9 2020 9 Alabama 188 26
## 10 2020 10 Alabama 183 30
## 11 2020 11 Alabama 217 51
## 12 2020 12 Alabama 546 153
## 13 2021 1 Alabama 838 175
## 14 2021 2 Alabama 334 65
## 15 2021 3 Alabama 138 32
## 16 2021 4 Alabama 70 15
## 17 2021 5 Alabama 74 18
## 18 2021 6 Alabama 43 NA
## 19 2021 7 Alabama 93 16
## 20 2021 8 Alabama 619 81
## 21 2021 9 Alabama 740 113
## 22 2021 10 Alabama 316 58
## 23 2021 11 Alabama 100 19
## 24 2021 12 Alabama 93 27
## 25 2022 1 Alabama 367 126
## # i 2,360 more rows
## # i 19 more variables: Adult.respiratory.distress.syndrome <int>,
## # Respiratory.failure <int>, Respiratory.arrest <int>,
## # Other.diseases.of.the.respiratory.system <int>,
## # Hypertensive.diseases <int>, Ischemic.heart.disease <int>,
## # Cardiac.arrest <int>, Cardiac.arrhythmia <int>, Heart.failure <int>,
## # Cerebrovascular.diseases <int>, ...
```

```
write_csv(wide_data, "wide_data_export.csv")
```

```
# We have lots of values with NA values
# We could assume some values, but that would introduce significant errors to the models.
# Instead, let's delete the rows that have NA values, this will reduce the rows in our dataset
# but there will be no errors introduced by the NA value due to counts of 1-9.
wide_data <- wide_data[(complete.cases(wide_data) == TRUE), ]
```

```

# Now lets also delete the rows that have 0 COVID-19 deaths.
# This was common at the beginning of 2020.
wide_data <- wide_data |> filter(`COVID19` > 0)

# Change the Month column to be a linear parameter as it gives us a better parameter for the models,
# no need for the Year parameter afterwards
# Experimented with using the month as a parameter as well
# But it did not provide much improvement, so eliminated the Month parameter later on..
wide_data <- wide_data |> mutate(Month = (Year - 2020)*12 + Month) |> select(-Year, -State, -Month)

#This is our final dataset
str(wide_data)

```

```

## tibble [430 x 21] (S3: tbl_df/tbl/data.frame)
##  $ Influenza.pneumonia          : int [1:430] 264 318 188 183 217 546 838 334 619 740 ...
##  $ Chronic.lower.respiratory.diseases : int [1:430] 73 55 26 30 51 153 175 65 81 113 ...
##  $ Adult.respiratory.distress.syndrome : int [1:430] 42 67 38 40 51 61 120 60 73 141 ...
##  $ Respiratory.failure          : int [1:430] 299 380 234 245 321 694 962 404 717 861 ...
##  $ Respiratory.arrest           : int [1:430] 74 71 37 39 62 158 196 86 175 193 ...
##  $ Other.diseases.of.the.respiratory.system: int [1:430] 26 34 22 22 26 71 90 54 51 75 ...
##  $ Hypertensive.diseases        : int [1:430] 117 105 55 51 60 170 228 108 145 179 ...
##  $ Ischemic.heart.disease       : int [1:430] 82 82 44 49 66 152 194 73 100 122 ...
##  $ Cardiac.arrest              : int [1:430] 116 128 61 68 115 287 370 171 298 348 ...
##  $ Cardiac.arrhythmia          : int [1:430] 60 58 31 33 64 122 181 59 106 117 ...
##  $ Heart.failure               : int [1:430] 58 55 35 38 50 128 165 59 78 95 ...
##  $ Cerebrovascular.diseases     : int [1:430] 30 41 21 27 27 72 106 45 53 58 ...
##  $ Other.diseases.of.the.circulatory.system: int [1:430] 42 61 40 35 51 108 181 66 101 132 ...
##  $ Sepsis                     : int [1:430] 87 112 70 68 89 174 284 139 153 240 ...
##  $ Malignant.neoplasms         : int [1:430] 35 12 24 22 25 66 90 39 49 42 ...
##  $ Diabetes                   : int [1:430] 95 110 68 58 79 147 251 111 153 187 ...
##  $ Obesity                    : int [1:430] 26 30 16 10 13 48 61 19 71 81 ...
##  $ Alzheimer.disease          : int [1:430] 44 31 16 17 14 40 41 18 13 16 ...
##  $ Vascular.unspecified.dementia : int [1:430] 68 71 65 39 38 94 108 42 30 33 ...
##  $ Renal.failure              : int [1:430] 96 82 57 48 92 176 281 110 167 216 ...
##  $ COVID19                   : int [1:430] 879 964 568 560 768 1749 2415 970 1682 1978

```

```

write_csv(wide_data, "wide_data_export.csv")

print('Mean of our output Parameter, COVID19:')

```

```
## [1] "Mean of our output Parameter, COVID19:"
```

```
mean(wide_data$COVID19)
```

```
## [1] 1828.651
```

```
print('Median of our output Parameter, COVID19:')

```

```
## [1] "Median of our output Parameter, COVID19:"
```

```
median(wide_data$COVID19)
```

```
## [1] 1236
```

Preparing training and test datasets

Let's prepare our training and test datasets

Final hold-out test set will be 10% of the data

We use the naming convention as in the previous project

```
# Set the seed for reproducibility  
set.seed(1, sample.kind="Rounding") # if using R 3.6 or later
```

```
## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding'  
## sampler used
```

```
# set.seed(1) # if using R 3.5 or earlier
```

```
test_index <- createDataPartition(y = wide_data$COVID19, times = 1, p = 0.1, list = FALSE)  
edx <- wide_data[-test_index,]  
final_holdout_test <- wide_data[test_index,]
```

Comparing the Models via Bootstrapping with a k value of 25

```
#####  
#  
# Comparing models with bootstrapping method with a value of 25  
# On my computer this subsection takes 2 minutes 50 seconds  
#  
  
# Define the trainControl with bootstrapping  
train_control <- trainControl(  
  method = "boot",      # Bootstrapping method  
  number = 25,          # Number of bootstrap samples  
  verboseIter = FALSE   # Set to TRUE to see the progress  
)  
  
# List of different model types to run  
models_list <- c("lm", "glm", "knn", "rf", "gamLoess", "rpart", "xgbTree",  
  "cforest", "glmnet", "bayesglm", "pcr", "pls", "ridge", "nnet")  
  
# Create empty lists to store the results  
boot_model_results <- list()  
boot_rmse <- list()  
  
# Loop over the list of models and train each one  
# suppress the warnings about deprecated functions
```

```

# and nnet training messages
for (model_type in models_list) invisible(capture.output({
  # to be able to replicate the results
  set.seed(1, sample.kind="Rounding") # if using R 3.6 or later
  # set.seed(1) # if using R 3.5 or earlier

  model <- train(
    COVID19 ~ .,          # COVID19 variable vs all predictors
    data = edx,           # Training dataset
    method = model_type,  # Current model from the list
    trControl = train_control # Use the defined trainControl
  )

  # Predict on the final_holdout_test dataset
  predictions <- predict(model, newdata = final_holdout_test)

  boot_rmse[[model_type]] <- sqrt(mean((predictions - final_holdout_test$COVID19)^2))

  # Store the result in the list
  boot_model_results[[model_type]] <- model
}))

```

```
## Loading required package: gam
```

```
## Loading required package: splines
```

```
## Loading required package: foreach
```

```
## Loaded gam 1.22-3
```

```
# Access results of each model
```

```
boot_rmse
```

```

## $lm
## [1] 291.7629
##
## $glm
## [1] 291.7629
##
## $knn
## [1] 252.6973
##
## $rf
## [1] 248.5894
##
## $gamLoess
## [1] 252.0772
##
## $rpart
## [1] 888.8269
##
## $xgbTree

```

```
## [1] 231.0553
##
## $cforest
## [1] 263.3392
##
## $glmnet
## [1] 304.0771
##
## $bayesglm
## [1] 291.7629
##
## $pcr
## [1] 342.2447
##
## $pls
## [1] 294.813
##
## $ridge
## [1] 291.7629
##
## $nnet
## [1] 2186.664
```

```
# Let's look at the best models and worst models
```

```
boot_top_3_rmse <- sort(unlist(boot_rmse), decreasing = FALSE)[1:3]
boot_top_3_rmse
```

```
## xgbTree      rf gamLoess
## 231.0553 248.5894 252.0772
```

```
boot_worst_3_rmse <- sort(unlist(boot_rmse), decreasing = TRUE)[1:3]
boot_worst_3_rmse
```

```
##      nnet      rpart      pcr
## 2186.6641 888.8269 342.2447
```

```
# Extract the variable importance for top 3 models
```

```
for (model_name in names(boot_top_3_rmse)) {
  cat("Variable Importance for model:", model_name, "\n")
  var_importance <- varImp(boot_model_results[[model_name]])
  print(var_importance)
  cat("\n")
}
```

```
## Variable Importance for model: xgbTree
```

```
## xgbTree variable importance
```

```
##
## Overall
## Influenza.pneumonia 100.00000
## Renal.failure 68.90390
## Respiratory.failure 45.86465
```

```

## Adult.respiratory.distress.syndrome      26.08815
## Respiratory.arrest                       16.18996
## Cardiac.arrhythmia                      14.16910
## Chronic.lower.respiratory.diseases       4.41823
## Ischemic.heart.disease                   4.28138
## Hypertensive.diseases                    2.35341
## Diabetes                                 2.32260
## Heart.failure                           1.18459
## Vascular.unspecified.dementia            0.77940
## Obesity                                 0.53307
## Cardiac.arrest                          0.38919
## Malignant.neoplasms                     0.38809
## Sepsis                                  0.22384
## Other.diseases.of.the.respiratory.system 0.16147
## Cerebrovascular.diseases                0.03492
## Other.diseases.of.the.circulatory.system 0.02575
## Alzheimer.disease                       0.00000
##
## Variable Importance for model: rf
## rf variable importance
##
## Overall
## Influenza.pneumonia                     100.0000
## Respiratory.failure                     84.0869
## Ischemic.heart.disease                   67.4816
## Cardiac.arrhythmia                      44.0161
## Other.diseases.of.the.circulatory.system 23.0288
## Cerebrovascular.diseases                16.9564
## Renal.failure                           16.5231
## Diabetes                                 13.6520
## Adult.respiratory.distress.syndrome     13.1184
## Respiratory.arrest                       9.0165
## Chronic.lower.respiratory.diseases       7.5493
## Heart.failure                           6.6621
## Hypertensive.diseases                    5.5820
## Cardiac.arrest                          4.9089
## Alzheimer.disease                       3.1151
## Malignant.neoplasms                     2.4882
## Other.diseases.of.the.respiratory.system 2.2472
## Sepsis                                  1.5487
## Vascular.unspecified.dementia            0.9601
## Obesity                                 0.0000
##
## Variable Importance for model: gamLoess
## gamLoess variable importance
##
## Overall
## Other.diseases.of.the.circulatory.system 100.000
## Cardiac.arrest                          84.522
## Renal.failure                           83.760
## Sepsis                                  80.355
## Cerebrovascular.diseases                68.651
## Chronic.lower.respiratory.diseases       57.188
## Diabetes                                 51.617

```

| | |
|---|--------|
| ## Obesity | 49.864 |
| ## Adult.respiratory.distress.syndrome | 49.663 |
| ## Respiratory.arrest | 45.461 |
| ## Hypertensive.diseases | 35.356 |
| ## Respiratory.failure | 34.682 |
| ## Other.diseases.of.the.respiratory.system | 27.160 |
| ## Influenza.pneumonia | 20.816 |
| ## Alzheimer.disease | 17.822 |
| ## Cardiac.arrhythmia | 16.083 |
| ## Heart.failure | 12.874 |
| ## Malignant.neoplasms | 8.244 |
| ## Vascular.unspecified.dementia | 3.550 |
| ## Ischemic.heart.disease | 0.000 |

Comparing the Models via 25-fold Cross-Validation

```
# Define the trainControl with cross validation method
train_control <- trainControl(
  method = "cv",          # cross validation method
  number = 25,            # Number of bootstrap samples
  verboseIter = FALSE     # Set to TRUE to see the progress
)

# Create empty lists to store the results
cv_model_results <- list()
cv_rmse <- list()

# Loop over the list of models and train each one
# suppress the warnings about deprecated functions
# and nnet training messages
for (model_type in models_list) invisible(capture.output({
  # to be able to replicate the results
  set.seed(1, sample.kind="Rounding") # if using R 3.6 or later
  # set.seed(1) # if using R 3.5 or earlier
  model <- train(
    COVID19 ~ .,          # COVID19 variable vs all predictors
    data = edx,           # Training dataset
    method = model_type,  # Current model from the list
    trControl = train_control # Use the defined trainControl
  )

  # Predict on the final_holdout_test dataset
  predictions <- predict(model, newdata = final_holdout_test)

  cv_rmse[[model_type]] <- sqrt(mean((predictions - final_holdout_test$COVID19)^2))

  # Store the result in the list
  cv_model_results[[model_type]] <- model
}))
```

```
# Access results of each model:
```

```
cv_rmse
```

```
## $lm
## [1] 291.7629
##
## $glm
## [1] 291.7629
##
## $knn
## [1] 252.6973
##
## $rf
## [1] 248.5894
##
## $gamLoess
## [1] 252.0772
##
## $rpart
## [1] 888.8269
##
## $xgbTree
## [1] 161.1663
##
## $cforest
## [1] 263.3392
##
## $glmnet
## [1] 304.0771
##
## $bayesglm
## [1] 291.7629
##
## $pcr
## [1] 342.2447
##
## $pls
## [1] 294.813
##
## $ridge
## [1] 292.1945
##
## $nnet
## [1] 2186.664
```

```
# Let's look at the best models and worst models
```

```
cv_top_3_rmse <- sort(unlist(cv_rmse), decreasing = FALSE)[1:3]
cv_top_3_rmse
```

```
## xgbTree      rf gamLoess
## 161.1663 248.5894 252.0772
```



```
cv_worst_3_rmse <- sort(unlist(cv_rmse), decreasing = TRUE)[1:3]
cv_worst_3_rmse
```

```
##      nnet      rpart      pcr
## 2186.6641  888.8269  342.2447
```

Extract the variable importance for top 3 models for 25-fold cross validation

```
# Extract the variable importance for top 3 models
for (model_name in names(cv_top_3_rmse)) {
  cat("Variable Importance for model:", model_name, "\n")
  var_importance <- varImp(cv_model_results[[model_name]])
  print(var_importance)
  cat("\n")
}
```

```
## Variable Importance for model: xgbTree
## xgbTree variable importance
##
##                                Overall
## Respiratory.failure            100.00000
## Influenza.pneumonia            90.07434
## Malignant.neoplasms            30.88395
## Adult.respiratory.distress.syndrome 21.18099
## Cardiac.arrhythmia             18.13848
## Respiratory.arrest             14.79050
## Chronic.lower.respiratory.diseases 12.53784
## Diabetes                       5.39707
## Ischemic.heart.disease          5.32439
## Heart.failure                   4.20999
## Hypertensive.diseases           2.35590
## Vascular.unspecified.dementia   1.06199
## Cardiac.arrest                  0.32535
## Sepsis                          0.29698
## Other.diseases.of.the.respiratory.system 0.22408
## Cerebrovascular.diseases        0.17067
## Renal.failure                   0.13228
## Obesity                         0.10758
## Alzheimer.disease               0.03441
## Other.diseases.of.the.circulatory.system 0.00000
##
## Variable Importance for model: rf
## rf variable importance
##
##                                Overall
## Influenza.pneumonia            100.0000
## Respiratory.failure             84.0869
## Ischemic.heart.disease          67.4816
## Cardiac.arrhythmia              44.0161
```

```

## Other.diseases.of.the.circulatory.system 23.0288
## Cerebrovascular.diseases 16.9564
## Renal.failure 16.5231
## Diabetes 13.6520
## Adult.respiratory.distress.syndrome 13.1184
## Respiratory.arrest 9.0165
## Chronic.lower.respiratory.diseases 7.5493
## Heart.failure 6.6621
## Hypertensive.diseases 5.5820
## Cardiac.arrest 4.9089
## Alzheimer.disease 3.1151
## Malignant.neoplasms 2.4882
## Other.diseases.of.the.respiratory.system 2.2472
## Sepsis 1.5487
## Vascular.unspecified.dementia 0.9601
## Obesity 0.0000
##
## Variable Importance for model: gamLoess
## gamLoess variable importance
##
## Overall
## Other.diseases.of.the.circulatory.system 100.000
## Cardiac.arrest 84.522
## Renal.failure 83.760
## Sepsis 80.355
## Cerebrovascular.diseases 68.651
## Chronic.lower.respiratory.diseases 57.188
## Diabetes 51.617
## Obesity 49.864
## Adult.respiratory.distress.syndrome 49.663
## Respiratory.arrest 45.461
## Hypertensive.diseases 35.356
## Respiratory.failure 34.682
## Other.diseases.of.the.respiratory.system 27.160
## Influenza.pneumonia 20.816
## Alzheimer.disease 17.822
## Cardiac.arrhythmia 16.083
## Heart.failure 12.874
## Malignant.neoplasms 8.244
## Vascular.unspecified.dementia 3.550
## Ischemic.heart.disease 0.000

```

The best models each way of training are

xgbTree,rf,gamLoess

The worst one is nnet with rmse value of 2186.6641

Is there a way to optimize the nnet model?

In this section, we change the number of units in the hidden layer and delay parameters

```
#####  
# Is there a way to optimize the nnet model?  
# In this section, we change the number of units in the hidden layer  
# and delay parameters  
# This section takes 2 minutes 55 seconds on my computer.  
  
# Define the control function with cross-validation  
train_control <- trainControl(method = "cv",  
                              number = 10,  
                              verboseIter = FALSE)  
  
# Define a grid of parameters to tune  
tune_grid <- expand.grid(  
  size = c(3, 4, 5, 6, 7, 8),    # Number of units in the hidden layer  
  decay = c(0.08, 0.09, 0.1, 0.11, 0.12, 0.13, 0.14, 0.15, 0.16)  
  # Regularization parameter to avoid overfitting  
)  
set.seed(1, sample.kind="Rounding") # if using R 3.6 or later  
# set.seed(1) # if using R 3.5 or earlier  
  
# Train the model using nnet and parameter tuning  
model <- train(  
  COVID19 ~ .,                # Formula: COVID19 variable vs all predictors  
  data = edx,  
  method = "nnet",  
  trControl = train_control,  
  tuneGrid = tune_grid,  
  linout = TRUE,              # For regression, use linout = TRUE  
  trace = FALSE,              # Avoid printing during training  
  maxit = 1000                # Maximum number of iterations  
)  
  
# View the best model  
print(model)
```

```
## Neural Network  
##  
## 386 samples  
## 20 predictor
```

```

##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 348, 347, 347, 347, 347, 347, ...
## Resampling results across tuning parameters:
##
##   size  decay  RMSE      Rsquared  MAE
##   3     0.08  1461.7288  0.4029400  900.9504
##   3     0.09  1304.5786  0.5561979  774.4989
##   3     0.10   892.4448  0.7081071  510.8347
##   3     0.11  1749.0244  0.4506128  973.2203
##   3     0.12  1420.1433  0.3894717  876.7805
##   3     0.13  1655.0017  0.4311037  941.2270
##   3     0.14  1104.6506  0.6282932  688.4912
##   3     0.15  1426.2666  0.4839861  857.9145
##   3     0.16  1451.3846  0.4050624  889.3228
##   4     0.08   970.6093  0.7991485  560.9270
##   4     0.09  1531.8160  0.4958368  932.6634
##   4     0.10  1549.9992  0.3944219  935.1142
##   4     0.11  1671.6249  0.3536426  949.9917
##   4     0.12  1129.2312  0.5827576  656.1703
##   4     0.13   991.0900  0.6130996  596.3633
##   4     0.14   867.6216  0.7692948  448.5631
##   4     0.15  1199.3275  0.5322934  758.9058
##   4     0.16  1297.4818  0.6228802  712.9032
##   5     0.08  1572.9826  0.4886247 1006.1169
##   5     0.09  1108.3497  0.7043281  642.1245
##   5     0.10  1461.2434  0.4473150  886.1084
##   5     0.11  1205.9990  0.5717576  764.2888
##   5     0.12  1311.1847  0.4309102  821.2130
##   5     0.13  1280.5665  0.5567895  697.3412
##   5     0.14  1158.1199  0.6538318  598.7541
##   5     0.15  1235.6017  0.5392565  755.1650
##   5     0.16  1621.5276  0.4033298  957.9995
##   6     0.08   961.2195  0.6477998  546.1922
##   6     0.09  1522.5762  0.4220780  944.0241
##   6     0.10  1665.0195  0.2887197 1005.9812
##   6     0.11   906.0025  0.6668108  579.8203
##   6     0.12  1350.3861  0.5198642  815.3678
##   6     0.13  1686.5648  0.2760743 1041.6823
##   6     0.14  1475.9919  0.4750050  859.4721
##   6     0.15  1234.6931  0.5149940  778.6927
##   6     0.16  1213.1139  0.6049648  653.6328
##   7     0.08  1455.2769  0.5317280  912.7069
##   7     0.09  1193.2547  0.5870303  769.2171
##   7     0.10  1122.7525  0.7330771  652.1454
##   7     0.11  1108.1631  0.6120605  669.3107
##   7     0.12  1554.5952  0.3402095  914.5665
##   7     0.13  1093.6465  0.6922732  636.5650
##   7     0.14  1298.6059  0.4531386  799.2176
##   7     0.15  1529.2012  0.4159417  949.1994
##   7     0.16  1149.9133  0.5122837  723.5033
##   8     0.08  1569.5018  0.4282105  960.9842
##   8     0.09  1352.8489  0.5861313  810.4373

```

```
##      8      0.10  1435.1006  0.4311863  874.9393
##      8      0.11  1556.8649  0.4417696  959.1880
##      8      0.12  1492.2099  0.3917677  959.4196
##      8      0.13  1287.9239  0.5348019  813.5884
##      8      0.14  1543.2916  0.4665072  860.6616
##      8      0.15  1351.4736  0.4433945  857.8633
##      8      0.16  1086.0209  0.5654945  658.5192
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were size = 4 and decay = 0.14.
```

```
print(model$bestTune)
```

```
##      size decay
## 16      4  0.14
```

```
model
```

```
## Neural Network
##
## 386 samples
## 20 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 348, 347, 347, 347, 347, ...
## Resampling results across tuning parameters:
##
##      size decay  RMSE      Rsquared  MAE
##      3      0.08 1461.7288  0.4029400  900.9504
##      3      0.09 1304.5786  0.5561979  774.4989
##      3      0.10  892.4448  0.7081071  510.8347
##      3      0.11 1749.0244  0.4506128  973.2203
##      3      0.12 1420.1433  0.3894717  876.7805
##      3      0.13 1655.0017  0.4311037  941.2270
##      3      0.14 1104.6506  0.6282932  688.4912
##      3      0.15 1426.2666  0.4839861  857.9145
##      3      0.16 1451.3846  0.4050624  889.3228
##      4      0.08  970.6093  0.7991485  560.9270
##      4      0.09 1531.8160  0.4958368  932.6634
##      4      0.10 1549.9992  0.3944219  935.1142
##      4      0.11 1671.6249  0.3536426  949.9917
##      4      0.12 1129.2312  0.5827576  656.1703
##      4      0.13  991.0900  0.6130996  596.3633
##      4      0.14  867.6216  0.7692948  448.5631
##      4      0.15 1199.3275  0.5322934  758.9058
##      4      0.16 1297.4818  0.6228802  712.9032
##      5      0.08 1572.9826  0.4886247 1006.1169
##      5      0.09 1108.3497  0.7043281  642.1245
##      5      0.10 1461.2434  0.4473150  886.1084
##      5      0.11 1205.9990  0.5717576  764.2888
##      5      0.12 1311.1847  0.4309102  821.2130
##      5      0.13 1280.5665  0.5567895  697.3412
```

```
## 5 0.14 1158.1199 0.6538318 598.7541
## 5 0.15 1235.6017 0.5392565 755.1650
## 5 0.16 1621.5276 0.4033298 957.9995
## 6 0.08 961.2195 0.6477998 546.1922
## 6 0.09 1522.5762 0.4220780 944.0241
## 6 0.10 1665.0195 0.2887197 1005.9812
## 6 0.11 906.0025 0.6668108 579.8203
## 6 0.12 1350.3861 0.5198642 815.3678
## 6 0.13 1686.5648 0.2760743 1041.6823
## 6 0.14 1475.9919 0.4750050 859.4721
## 6 0.15 1234.6931 0.5149940 778.6927
## 6 0.16 1213.1139 0.6049648 653.6328
## 7 0.08 1455.2769 0.5317280 912.7069
## 7 0.09 1193.2547 0.5870303 769.2171
## 7 0.10 1122.7525 0.7330771 652.1454
## 7 0.11 1108.1631 0.6120605 669.3107
## 7 0.12 1554.5952 0.3402095 914.5665
## 7 0.13 1093.6465 0.6922732 636.5650
## 7 0.14 1298.6059 0.4531386 799.2176
## 7 0.15 1529.2012 0.4159417 949.1994
## 7 0.16 1149.9133 0.5122837 723.5033
## 8 0.08 1569.5018 0.4282105 960.9842
## 8 0.09 1352.8489 0.5861313 810.4373
## 8 0.10 1435.1006 0.4311863 874.9393
## 8 0.11 1556.8649 0.4417696 959.1880
## 8 0.12 1492.2099 0.3917677 959.4196
## 8 0.13 1287.9239 0.5348019 813.5884
## 8 0.14 1543.2916 0.4665072 860.6616
## 8 0.15 1351.4736 0.4433945 857.8633
## 8 0.16 1086.0209 0.5654945 658.5192
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were size = 4 and decay = 0.14.
```

```
# Predict on the final_holdout_test dataset
predictions <- predict(model$finalModel, newdata = final_holdout_test)

# improved RMSE:
sqrt(mean((predictions - final_holdout_test$COVID19)^2))
```

```
## [1] 1492.863
```

```
# RMSE has improved significantly, 1492.863
# but still it is much worse than our best model xgbTree: 161
# apparently nnet can only have one layer of hidden layer.
# To have multiple hidden layers, keras package from Google can be used apparently.
```